

```
In [1]: import cv2
import os
import random
import numpy as np
from matplotlib import pyplot as plt
import uuid
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Layer, Conv2D, Dense, MaxPooling2D, Input
import tensorflow as tf
from tensorflow.keras.metrics import Precision, Recall
```

```
In [2]: cwd = os.getcwd()
```

```
In [3]: cwd
```

```
Out[3]: '/home/kumar.amit1/ondemand/data/sys/myjobs/projects/default/1/data'
```

```
In [4]: POS_PATH = 'data_crop/positive_crop'
NEG_PATH = 'data_crop/negative_crop'
ANC_PATH = 'data_crop/anchor_crop'
```

Loading and Processing Image

```
In [5]: # Get Image directory
anchor = tf.data.Dataset.list_files(ANC_PATH+'/*.jpg').take(3000)
positive = tf.data.Dataset.list_files(POS_PATH+'/*.jpg').take(3000)
negative = tf.data.Dataset.list_files(NEG_PATH+'/*.jpg').take(3000)
```

2023-04-25 22:24:44.706003: I tensorflow/core/platform/cpu_feature_guard.cc:142] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SSE4.1 SSE4.2 AVX AVX2 FMA

To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

2023-04-25 22:24:45.519464: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1510] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 77292 MB memory: -> device: 0, name: NVIDIA A100-SXM4-80GB, pci bus id: 0000:bd:00.0, compute capability: 8.0

```
In [6]: # Scale and Resize
def preprocess(file_path):

    # Read in image from file path
    byte_img = tf.io.read_file(file_path)
    # Load in the image
    img = tf.io.decode_jpeg(byte_img)

    # Preprocessing steps - resizing the image to be 100x100x3
    img = tf.image.resize(img, (300,300))
    # Scale image to be between 0 and 1
    img = img / 255.0

    # Return image
    return img
```

Creating Labels

```
In [7]: positives = tf.data.Dataset.zip((anchor, positive, tf.data.Dataset.from_tensor_slices(
negatives = tf.data.Dataset.zip((anchor, negative, tf.data.Dataset.from_tensor_slices(
data = positives.concatenate(negatives)
```

```
In [8]: samples = data.as_numpy_iterator()
```

```
In [9]: example = samples.next()
print(example)
```


```
(b'data_crop/anchor_crop/RE2222.jpg', b'data_crop/positive_crop/RE326.jpg',
1.0)
```

Train and Test Partiton

```
In [10]: def preprocess_twain(input_img, validation_img, label):
return(preprocess(input_img), preprocess(validation_img), label)
```

```
In [11]: res = preprocess_twin(*example)
plt.imshow(res[1])
print(res[2])
```

1.0



```
In [12]: # Build dataloader pipeline
data = data.map(preprocess_twin)
data = data.cache()
data = data.shuffle(buffer_size=10000)
```

```
In [13]: # Training partition
train_data = data.take(round(len(data)*.7))
train_data = train_data.batch(16)
train_data = train_data.prefetch(8)
```

```
In [14]: #Testing partition
test_data = data.skip(round(len(data)*.7))
test_data = test_data.take(round(len(data)*.3))
test_data = test_data.batch(16)
test_data = test_data.prefetch(8)
```

Model Engineering

Embedding Model

```
In [15]: def form_embedding():
    #Input Layer
    inputlayer = Input(shape=(300,300,3), name='input_image')

    #First Layer
    c1 = Conv2D(64, (10,10), activation='relu')(inputlayer)
    m1 = MaxPooling2D(64, (2,2), padding='same')(c1)

    # Second Layer
    c2 = Conv2D(128, (7,7), activation='relu')(m1)
    m2 = MaxPooling2D(64, (2,2), padding='same')(c2)

    # Third Layer
    c3 = Conv2D(128, (4,4), activation='relu')(m2)
    m3 = MaxPooling2D(64, (2,2), padding='same')(c3)

    #Final Layer
    c4 = Conv2D(256, (4,4), activation='relu')(m3)
    f1 = Flatten()(c4)
    d1 = Dense(4096, activation='sigmoid')(f1)

    return Model(inputs=[inputlayer], outputs=[d1], name='embedding')
```

```
In [16]: embedding=form_embedding()
```

distance Layer

```
In [17]: # Siamese L1 Distance class
class L1Dist(Layer):

    # Init method - inheritance
    def __init__(self, **kwargs):
        super().__init__()

    # Magic happens here - similarity calculation
    def call(self, input_embedding, validation_embedding):
        return tf.math.abs(input_embedding - validation_embedding)
```

Siamese network

```
In [18]: def make_siamese_model():

    # Anchor image input in the network
    input_image = Input(name='input_img', shape=(300,300,3))

    # Validation image in the network
    validation_image = Input(name='validation_img', shape=(300,300,3))

    # Combine siamese distance components
    siamese_layer = L1Dist()
    siamese_layer._name = 'distance'
    #distances = siamese_layer(embedding_efB4(input_image), embedding_efB4(va
    distances = siamese_layer(embedding(input_image), embedding(validation_im

    # Classification layer
    classifier = Dense(1, activation='sigmoid')(distances)

    return Model(inputs=[input_image, validation_image], outputs=classifier,

In [19]: siamese_model = make_siamese_model()
```

```
In [20]: siamese_model.summary()
```

Model: "SiameseNetwork"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_img (InputLayer)	[(None, 300, 300, 3)]	0	
=====			
validation_img (InputLayer)	[(None, 300, 300, 3)]	0	
=====			
embedding (Functional) [0][0]	(None, 4096)	1008893248	input_img validation_ img[0][0]
=====			
distance (L1Dist) [0][0] [1][0]	(None, 4096)	0	embedding embedding
=====			
dense_1 (Dense) [0]	(None, 1)	4097	distance[0]
=====			
Total params: 1,008,897,345			
Trainable params: 1,008,897,345			
Non-trainable params: 0			

Training

```
In [21]: # Loss and Optimier
binary_cross_loss = tf.losses.BinaryCrossentropy()
opt = tf.keras.optimizers.Adam(1e-4)
```

```
In [22]: # establish checkpoint
checkpoint_dir = './training_checkpoints'
checkpoint_prefix = os.path.join(checkpoint_dir, 'ckpt')
checkpoint = tf.train.Checkpoint(opt=opt, siamese_model=siamese_model)
```

```

In [23]: # Train Step Function
@tf.function
def train_step(batch):

    # Record all of our operations
    with tf.GradientTape() as tape:
        # Get anchor and positive/negative image
        X = batch[:2]
        # Get Label
        y = batch[2]

        # Forward pass
        yhat = siamese_model(X, training=True)
        # Calculate Loss
        loss = binary_cross_loss(y, yhat)
    print(loss)

    # Calculate gradients
    grad = tape.gradient(loss, siamese_model.trainable_variables)

    # Calculate updated weights and apply to siamese model
    opt.apply_gradients(zip(grad, siamese_model.trainable_variables))

    # Return Loss
    return loss

```

```

In [24]: # Training Loop
def train(data, EPOCHS):
    # Loop through epochs
    for epoch in range(1, EPOCHS+1):
        print('\n Epoch {}/{}'.format(epoch, EPOCHS))
        progbar = tf.keras.utils.Progbar(len(data))

        # Creating a metric object
        r = Recall()
        p = Precision()

        # Loop through each batch
        for idx, batch in enumerate(data):
            # Run train step here
            loss = train_step(batch)
            yhat = siamese_model.predict(batch[:2])
            r.update_state(batch[2], yhat)
            p.update_state(batch[2], yhat)
            progbar.update(idx+1)
        print(loss.numpy(), r.result().numpy(), p.result().numpy())

    # Save checkpoints
    if epoch % 10 == 0:
        checkpoint.save(file_prefix=checkpoint_prefix)

```

Train The model

```
In [25]: EPOCHS = 5
```

```
In [26]: train(train_data, EPOCHS)
```

Epoch 1/5

```
2023-04-25 22:24:46.997862: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the MLIR Optimization Passes are enabled (registered 2)
```

```
2023-04-25 22:24:56.998202: I tensorflow/core/kernels/data/shuffle_dataset_op.cc:175] Filling up shuffle buffer (this may take a while): 785 of 10000
```

```
2023-04-25 22:24:57.633292: I tensorflow/core/kernels/data/shuffle_dataset_op.cc:228] Shuffle buffer filled.
```

```
Tensor("binary_crossentropy/weighted_loss/value:0", shape=(), dtype=float32)
Tensor("binary_crossentropy/weighted_loss/value:0", shape=(), dtype=float32)
```

```
2023-04-25 22:24:58.944432: I tensorflow/stream_executor/cuda/cuda_dnn.cc:369] Loaded cuDNN version 8201
```

```
2023-04-25 22:24:59.986253: W tensorflow/stream_executor/gpu/asm_compiler.cc:77] Couldn't get ptxas version string: Internal: Running ptxas --version returned 32512
```

```
2023-04-25 22:25:00.074473: W tensorflow/stream_executor/gpu/redzone_allocator.cc:314] Internal: ptxas exited with non-zero error code 32512, output: Relying on driver to perform ptx compilation.
```

```
Modify $PATH to customize ptxas location.
```

```
This message will be only logged once.
```

```
2023-04-25 22:25:01.890914: I tensorflow/stream_executor/cuda/cuda_blas.cc:1760] TensorFlow-32 will be used for the matrix multiplication. This will only be logged once.
```

```
44/45 [=====>.] - ETA: 0sTensor("binary_crossentropy/weighted_loss/value:0", shape=(), dtype=float32)
```

```
45/45 [=====] - 48s 665ms/step
```

```
0.021557642 0.9854651 0.96033996
```

Epoch 2/5

```
45/45 [=====] - 29s 634ms/step
```

```
0.043293778 1.0 1.0
```

Epoch 3/5

```
45/45 [=====] - 29s 633ms/step
```

```
1.2546935e-05 0.9971671 0.9971671
```

Epoch 4/5

```
45/45 [=====] - 28s 627ms/step
```

```
0.0017192084 0.99712646 0.9942693
```

Epoch 5/5

```
45/45 [=====] - 28s 615ms/step
```

```
0.0038640506 0.9941691 0.98554915
```


Evaluating Model

```
In [27]: # Batch of test data
test_input, test_val, y_true = test_data.as_numpy_iterator().next()
y_hat = siamese_model.predict([test_input, test_val])
```

```
In [28]: # Post processing the results
[1 if prediction > 0.5 else 0 for prediction in y_hat ]
print(y_true)
```

```
[1. 1. 0. 0. 0. 1. 1. 1. 0. 1. 0. 0. 1. 1. 1. 1.]
```

Comparision Matrix

```
In [29]: # Creating a metric object
m = Recall()

# Calculating the recall value
m.update_state(y_true, y_hat)

# Return Recall Result
m.result().numpy()
```

```
Out[29]: 1.0
```

```
In [30]: r = Recall()
p = Precision()

for test_input, test_val, y_true in test_data.as_numpy_iterator():
    yhat = siamese_model.predict([test_input, test_val])
    r.update_state(y_true, yhat)
    p.update_state(y_true, yhat)

print(r.result().numpy(), p.result().numpy())
```

```
1.0 1.0
```


Result

```
In [31]: # Set plot size
plt.figure(figsize=(10,8))

# Set first subplot
plt.subplot(1,2,1)
plt.imshow(test_input[5])

# Set second subplot
plt.subplot(1,2,2)
plt.imshow(test_val[5])

# Renders cleanly
plt.show()
```



Save Model

```
In [ ]: #Save model weights
#siamese_model.save('siamesemodelv2.h5')
```

WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

2023-04-25 22:27:33.489992: W tensorflow/python/util/util.cc:348] Sets are not currently considered sequences, but this may change in the future, so consider avoiding using them.

```
In [ ]: # Reload model
#siamese_model_predict = tf.keras.models.load_model('siamesemodelv2.h5', custom_objects={'SiameseModel': SiameseModel})
```

```
In [ ]: # Make predictions with reloaded model  
        #siamese_model.predict([test_input, test_val])
```

```
In [ ]:
```